

Location Aware Keyword Query Suggestion Based on Document Proximity

Shuyao Qi, Dingming Wu, and Nikos Mamoulis

Abstract—Keyword suggestion in web search helps users to access relevant information without having to know how to precisely express their queries. Existing keyword suggestion techniques do not consider the locations of the users and the query results; i.e., the spatial proximity of a user to the retrieved results is not taken as a factor in the recommendation. However, the relevance of search results in many applications (e.g., location-based services) is known to be correlated with their spatial proximity to the query issuer. In this paper, we design a location-aware keyword query suggestion framework. We propose a weighted keyword-document graph, which captures both the semantic relevance between keyword queries and the spatial distance between the resulting documents and the user location. The graph is browsed in a random-walk-with-restart fashion, to select the keyword queries with the highest scores as suggestions. To make our framework scalable, we propose a partition-based approach that outperforms the baseline algorithm by up to an order of magnitude. The appropriateness of our framework and the performance of the algorithms are evaluated using real data.

Index Terms—Query suggestion, spatial databases

1 INTRODUCTION

KEYWORD suggestion (also known as query suggestion) has become one of the most fundamental features of commercial web search engines. After submitting a keyword query, the user may not be satisfied with the results, so the keyword suggestion module of the search engine recommends a set of m keyword queries that are most likely to refine the user's search in the right direction. Effective keyword suggestion methods are based on click information from query logs [1], [2], [3], [4], [5], [6], [7] and query session data [8], [9], [10], or query topic models [11]. New keyword suggestions can be determined according to their semantic relevance to the original keyword query.

However, to our knowledge, none of the existing methods provide *location-aware* keyword query suggestion (LKS), such that the suggested queries retrieve documents not only related to the user information needs but also located near the user location. This requirement emerges due to the popularity of spatial keyword search [12], [13], [14], [15], [16]. Google processed a daily average of 4.7 billion queries in 2011,¹ a substantial fraction of which have local intent and target spatial web objects (i.e., points of interest with a web presence having locations as well as text

descriptions) or geo-documents (i.e., documents associated with geo-locations). Furthermore, 53 percent of Bing's mobile searches in 2011 have a local intent.²

In this paper, we propose the first Location-aware Keyword query Suggestion framework. We illustrate the benefit of LKS using a toy example. Consider five geo-documents d_1-d_5 as listed in Fig. 1a. Each document d_i is associated with a location $d_i.\lambda$ as shown in Fig. 1b. Assume that a user issues a keyword query $k_q = "seafood"$ at location λ_q , shown in Fig. 1b. Note that the relevant documents d_1-d_3 (containing "seafood") are far from λ_q . A location-aware suggestion is "lobster", which can retrieve nearby documents d_4 and d_5 that are also relevant to the user's original search intention. Previous keyword query suggestion models (e.g., [5]) ignore the user location and would suggest "fish", which again fails to retrieve nearby relevant documents. Note that LKS has a different goal and therefore differs from other location-aware recommendation methods (e.g., auto-completion/instant search [17], [18], tag recommendation [19]). Section 5 provides a detailed discussion about the differences between LKS and these models, while in Section 4 we experimentally show that an adaptation of the method in [19] is less effective than LKS.

The first challenge of our LKS framework is how to effectively measure keyword query similarity while capturing the spatial distance factor. In accordance to previous query suggestion approaches [2], [3], [4], [5], [6], [7], [9], [10], LKS constructs and uses a keyword-document bipartite graph (KD-graph for short), which connects the keyword queries with their relevant documents as shown in Fig. 1c. Different to all previous approaches which ignore locations, LKS adjusts the weights on edges in the KD-graph to capture not only the semantic relevance between keyword queries, but also the spatial distance between the document locations and the query issuer's location λ_q . We apply the popular graph distance measure, *Personalized PageRank* (PPR) a.k.a.

1. <http://www.statisticbrain.com/google-searches>

• S. Qi is with the Department of Computer Science, the University of Hong Kong, Hong Kong. E-mail: qisy@connect.hku.hk.
 • N. Mamoulis is with the Department of Computer Science and Engineering, University of Ioannina. E-mail: nikos@cs.uoi.gr.
 • D. Wu is with the College of Computer Science & Software Engineering, Shenzhen University, China, and the Department of Computer Science, University of Hong Kong, Hong Kong. E-mail: dmwu@cs.hku.hk.

Manuscript received 12 Mar. 2015; revised 31 July 2015; accepted 1 Aug. 2015. Date of publication 6 Aug. 2015; date of current version 3 Dec. 2015. Recommended for acceptance by W. Wang.

For information on obtaining reprints of this article, please send e-mail to: reprints@ieee.org, and reference the Digital Object Identifier below.

Digital Object Identifier no. 10.1109/TKDE.2015.2465391

Authorized licensed use limited to: SHENZHEN UNIVERSITY. Downloaded on January 28, 2026 at 00:53:37 UTC from IEEE Xplore. Restrictions apply.

1041-4347 © 2015 IEEE. Personal use is permitted, but republication/redistribution requires IEEE permission.

See http://www.ieee.org/publications_standards/publications/rights/index.html for more information.

2. <http://searchengineland.com/microsoft-53-percent-of-mobile-searches-have-local-intent-55556>




Fig. 1. LKS example.

random walk with restart (RWR) distance [20] on the KD-graph, starting from the user supplied query k_q , to find the set of m keyword queries with the highest semantic relevance to k_q and spatial proximity to the user location λ_q . RWR on a KD-graph has been considered superior to alternative approaches [6] and has been a standard technique employed in previous (location-independent) keyword query suggestion studies [4], [5], [6], [7], [9], [10].

The second challenge is how to compute the suggestions efficiently. Performing keyword suggestion instantly is important for the applicability of LKS in practice. However, RWR search has a high computational cost on large graphs. Previous work on scaling up RWR search require pre-computation and/or graph segmentation [20], [21], [22], [23], [24] and require that the transition probabilities between nodes (i.e., the edge weights) are known beforehand. However, the edge weights of our KD-graph are unknown in advance, hindering the application of all these approaches. To the best of our knowledge, no existing technique can accelerate RWR when edge weights are unknown apriori (or they are dynamic). To address this issue, we present a novel partition-based algorithm (PA) that greatly reduces the cost of RWR search on such a dynamic bipartite graph. In a nutshell, our proposal divides the keyword queries and the documents into partitions and adopts a lazy mechanism that accelerates RWR search. PA and the lazy mechanism are generic techniques for RWR search, orthogonal to LKS, therefore they can be applied to speed up RWR search in other large graphs.

In summary, the contributions of this paper are:

- We design the first ever Location-aware Keyword query Suggestion framework, for suggestions relevant to the user's information needs that also retrieve relevant documents close to the query issuer's location.
- We extend the state-of-the-art Bookmark Coloring Algorithm (BCA) [25] for RWR search to compute the location-aware suggestions. In addition, we propose a partition-based algorithm (PA) that greatly reduces the computational cost of BCA.
- We conduct an empirical study that demonstrates the usefulness of location-aware keyword query suggestion. We also show experimentally that PA is two times to one order of magnitude faster than BCA.

The rest of the paper is organized as follows. LKS is introduced in Section 2. Our partition-based algorithm is presented in Section 3. We evaluate the effectiveness of LKS

and the performance of PA in Section 4. Related work is reviewed in Section 5 and we conclude in Section 6.

2 LKS FRAMEWORK

Consider a user-supplied query q with initial input k_q ; k_q can be a single word or a phrase. Assuming that the query issuer is at location λ_q , two intuitive criteria for selecting good suggestions are: (i) the suggested keyword queries (words or phrases) should satisfy the user's information needs based on k_q and (ii) the suggested queries can retrieve relevant documents spatially close to λ_q . The proposed LKS framework captures these two criteria.

2.1 Initial Keyword-Document Graph

Without loss of generality, we consider a set of geo-documents D such that each document $d_i \in D$ has a point location $d_i \cdot \lambda$.³ Let K be a collection of keyword queries from a query log. LKS first constructs an *initial* keyword-document graph (KD-graph), which is what a classic keyword suggestion approach that does not consider locations would use [4], [5], [6], [7], [9], [10]. This directed weighted bipartite graph $G = (D, K, E)$ between D and K captures the semantics and textual relevance between the keyword query and document nodes; i.e., the first criterion of location-aware suggestion. If a document d_i is clicked by a user who issued keyword query k_j in the query log, E contains an edge e from k_j to d_i and an edge e' from d_i to k_j . The weights of edges e and e' are the same and equal to the number of clicks on document d_i , given keyword query k_j [1]. Therefore, the direct relevance between a keyword query and a clicked document is captured by the edge weight. Furthermore, the semantic relevance between two keyword queries is captured by their proximity in the graph G (e.g., computed as their RWR distance). Any updates in the query log and/or the document database can be easily applied on the KD-graph; for a new query/document, we add a new node to the graph; for new clicks, we only need to update the corresponding edge weights accordingly. As an example, Fig. 1a shows five documents d_1-d_5 and three keyword queries k_1-k_3 . The corresponding KD-graph is shown in Fig. 1c. For the ease of presentation, the edge weights are normalized (i.e., divided by the maximum number of clicks in the log for any query-document pair).

³ If a document relates to multiple locations, we can model it as multiple documents, each referring to a single location. Location-independent documents can also be included in our framework by turning off the location awareness component for them.

2.2 Location-Aware Edge Weight Adjustment

In order to satisfy the second criterion of location-aware suggestion (i.e., location awareness), we propose to *adjust* the edge weights in the KD-graph based on the spatial relationships between the location of the query issuer and the nodes of the KD-graph. Note that this edge adjustment is query-dependent and dynamic. In other words, different adjustment is used for each different query independently.

We now outline the details of the edge weights adjustment. Recall that a user-supplied query q consists of two arguments: an input keyword query k_q (a word or a phrase) and a query location λ_q . Given q , the weight $w(e)$ of the edge e from a keyword query node k_i to a document node d_j is adjusted by the following function:

$$\tilde{w}(e) = \beta \times w(e) + (1 - \beta) \times (1 - dist(\lambda_q, d_j, \lambda)), \quad (1)$$

where $w(e)$ is the initial weight of e in the KD-graph, $\tilde{w}(e)$ is the adjusted edge weight, $dist(\lambda_q, d_j, \lambda)$ is the euclidean distance between the query issuer's location λ_q and document d_j , and parameter $\beta \in [0, 1]$ is used to balance the importance between the original (i.e., click-based) weight and the distance of d_j to the query location. Euclidean distances are normalized to take values in $[0, 1]$. This keyword-to-document edge weight adjustment increases the weights of the documents that are close to the user's location.

Let $D(k_i)$ be the set of documents connected to a keyword query $k_i \in K$ in the KD-graph. $D(k_i)$ may contain multiple documents and the locations of them form a spatial distribution. We propose to adjust the weights of the edges pointing to k_i by the minimum distance between λ_q and the locations of documents in $D(k_i)$.⁴ Such an adjustment favors keyword query nodes which have at least one relevant document close to the query issuer's location λ_q . Specifically, the weight $w(e')$ of the edge e' from a document node d_j to a keyword query node k_i is adjusted as follows:

$$\tilde{w}(e') = \beta \times w(e') + (1 - \beta) \times (1 - mindist(\lambda_q, D(k_i))), \quad (2)$$


where $mindist(\lambda_q, D(k_i))$ is the minimum euclidean distance⁵ between λ_q and any document in $D(k_i)$.

For example, Fig. 1b shows the locations of the five documents of Fig. 1a and a query location λ_q ; Fig. 1d includes the (approximate) euclidean distances between λ_q and the five documents. Fig. 2 illustrates how the edge weights from keyword query nodes to document nodes (Fig. 2a) and from document nodes to keyword query nodes (Fig. 2b) are adjusted based on the query location, assuming $\beta = 0.5$. Take the edge from k_1 to d_1 as a concrete example. Its weight is calculated using Equation (1) where $dist(\lambda_q, d_1, \lambda) = 1$. The weight of the edge from d_1 to k_1 is computed using Equation (2) where $D(k_1) = \{d_1, d_2\}$ and $mindist(\lambda_q, D(k_1)) = 0.9$.

We remark that the original KD-graph G is constructed only once in advance (as in previous work [4], [5], [6], [7],

4. Since the locations of past query issuers are not always available (e.g., due to privacy constraints), in this paper, we focus on the case where only document locations are known. Therefore, the edge adjustments for keyword-to-document edges and document-to-keyword edges are performed differently.

5. The effect of using the average distance to $D(k_i)$ is similar.



(a) Keyword queries → Documents

(b) Keyword queries ← Documents

Fig. 2. Location-aware edge weight adjustment.

[9], [10]). In addition, any update operations on the KD-graph (discussed in Section 2.1) are independent to our edge weight adjustment strategy, which is query-dependent. Given a user-supplied query q , the adjusted graph G_q is *dynamically* derived from G based on the query location λ_q , used to compute suggestions for q , and then dropped. During this process, G_q is maintained separately and G is not changed, so that concurrent or follow up queries are not affected. As we will discuss in Section 3.1, only a small portion of edges, relevant to the current query, are adjusted and cached, hence the adjustment is conducted efficiently and on-demand, during the keyword query suggestion process.

2.3 Location-Aware Keyword Query Suggestion

We denote by G_q the KD-graph G after adjusting the edge weights, based on the query location λ_q . G_q captures the two criteria of selecting suggestions, i.e., relevance to k_q and closeness to λ_q . Thus, keyword queries close to k_q in G_q are likely to be relevant to k_q and, at the same time, they result in documents close to the query issuer. In order to find the set of keyword queries for recommendation, we compute for all keyword queries a graph proximity score with respect to k_q , based on the random walk with restart process (typically used to measure graph proximity). Intuitively, the RWR score of a node v in graph G_q models the probability that a random surfer starting from k_q will reach v . At each step of the walk, the surfer either moves to an adjacent node with a probability $1 - \alpha$ (the next node depends on the weight of the corresponding edge), or ‘teleports’ to k_q with a probability α . The top- m keyword nodes in G_q with the highest scores (excluding k_q) are the suggestions.

Formally, let $\vec{\psi}$ be a column vector recording the RWR scores of all keyword queries in K based on G_q . $\vec{\psi}$ is computed by [26]:

$$\vec{\psi} = (1 - \alpha) M_{DK}^T M_{KD}^T \vec{\psi} + \alpha \vec{\psi}_q. \quad (3)$$

M_{DK} is a document-by-keyword matrix and M_{KD} is a keyword-by-document matrix, storing the edge weights in G_q ; both matrices are row-normalized. $\vec{\psi}_q$ is the initial score vector having zeros at all positions except the position of k_q , where it has 1. Since the user-supplied query k_q also gets an RWR score, in the end we compute the top- m keyword queries other than k_q .

Example 1. Consider query q with initial input $k_q = "seafood"$ and location $\lambda_q = (0.2, 0.2)$ against the




Fig. 3. Illustration of Algorithm BA.

documents of Fig. 1a. The locations of the documents and λ_q are shown in Fig. 1b. The KD-graph G is shown in Fig. 1c and the adjusted edge weights based on λ_q are presented in Fig. 2. The top-1 suggestion for q is “lobster” according to Equation (3) given $\alpha = 0.5$. Note that “lobster” retrieves documents d_4 and d_5 that are close to the query location, which cannot be achieved by “seafood”. Also queries “lobster” and “seafood” are semantically relevant.

3 ALGORITHMS

In this section, we introduce a baseline algorithm (BA) for location-aware suggestions (Section 3.1). Then, we propose our efficient partition-based algorithm (Section 3.2).

3.1 Baseline Algorithm (BA)

We extend the popular Bookmark-Coloring Algorithm [25] to compute the RWR-based top- m query suggestions as a baseline algorithm. BCA models RWR as a *bookmark coloring* process. Starting with one unit of active *ink* injected into node k_q , BA processes the nodes in the graph in descending order of their active ink. Different from typical personalized PageRank problems [27], [28] where the graph is homogeneous, our KD-graph G_q has two types of nodes: keyword query nodes and document nodes. As opposed to BCA, BA only ranks keyword query nodes; a keyword query node *retains* α portion of its active ink and *distributes* $1 - \alpha$ portion to its neighbor nodes based on its outgoing adjusted edge weights, while a document node distributes *all* its active ink to its neighbor nodes.

In our implementation, the weight of each edge e is adjusted based on λ_q *online*, at the time when the source node of e is distributing ink. This means that the edge weight adjustment which we propose in Section 2.2 is done during BA (i.e., G_q needs not be computed and materialized before BA starts). Moreover, a node may be processed several times; thus, the adjusted weights of its outgoing edges are cached after the node is first processed, for later usage. A node can distribute ink when its active ink exceeds a threshold ϵ . Algorithm BA terminates when either (i) the ink retained at the top- m th keyword query node is more than the ink retained at the top- $(m+1)$ th keyword query node plus the sum of the active ink of all nodes [27] or (ii) the active ink of each node is less than ϵ (typically, $\epsilon = 10^{-5}$).

Algorithm 1 is a pseudo code of BA. Priority queue Q maintains the nodes to be processed in descending order of their active ink (line 1). Q initially contains one entry, i.e., the user-supplied keywords k_q with active ink 1 (line 2). Priority

queue C , initially empty, stores the candidate suggestions in descending order of their retained ink (line 1). The sum of the active ink of all nodes $AINK$ is set to 1 (line 3). Termination conditions (i) and (ii) are checked at lines 4 and 8, respectively. The processing of a keyword query node involves retaining α portion of its active ink (line 13) and distributing $1 - \alpha$ portion to its neighbor document nodes based on the adjusted edge weights (lines 19-23). The total active ink $AINK$ is modified accordingly (line 14). As soon as a keyword query node has some retained ink, it enters C . The processing of a document node involves distributing all its active ink to neighbor keyword query nodes according to the adjusted edge weights (lines 19-23). The algorithm returns the top- m candidate suggestions other than k_q in C as the result (line 24).

Algorithm 1. Baseline Algorithm (BA)

```

Input:  $G(D, K, E)$ ,  $q = (k_q, \lambda_q)$ ,  $m, \epsilon$ 
Output:  $C$ 
1 PriorityQueue  $Q \leftarrow \emptyset$ ,  $C \leftarrow \emptyset$ ;
2 Add  $k_q$  to  $Q$  with  $k_q.aink \leftarrow 1$ ;
3  $AINK \leftarrow 1$ ;
4 while  $Q \neq \emptyset$  and  $Q.top.aink \geq \epsilon$  do
5   Deheap the first entry  $top$  from  $Q$ ;
6    $tm =$  the top- $m$  entry from  $C$ ;
7    $tm' =$  the top- $(m+1)$  entry from  $C$ ;
8   if  $tm.rank > tm'.rank + AINK$  then
9     break
10     $distratio = 1$  ;
11    if  $top$  is a keyword query node then
12       $distratio = 1 - \alpha$  ;
13       $top.rank \leftarrow top.rank + top.aink \times \alpha$ ;
14       $AINK \leftarrow AINK - top.aink \times \alpha$ ;
15      if there exist a copy  $t$  of  $top$  in  $C$  then
16        Remove  $t$  from  $C$ ;
17         $top.rank \leftarrow top.rank + t.rank$ ;
18        Add  $top$  to  $C$ ;
19    for each node  $v$  connected to  $top$  in  $G$  do
20       $v.aink \leftarrow top.aink \times distratio \times \tilde{w}(top, v)$ ;
21      if there exists a copy  $v'$  of  $v$  in  $Q$  then
22        Remove  $v'$  from  $Q$ ;  $v.aink \leftarrow v.aink + v'.aink$ ;
23        Add  $v$  to  $Q$ ;
24 return the top- $m$  entries (excluding  $k_q$ ) in  $C$ ;

```

Example 2. Fig. 3 shows the steps of BA (for $m = 1$, $\epsilon = 0.1$ and $\alpha = 0.5$), when applied to the adjusted KD graph of our running example (see Example 1 and Figs. 1 and 2). The number next to each node indicates its amount of active ink. The numbers in rounded rectangles are the amount of retained ink. Initially, one unit of ink is injected into node k_2 , i.e., the keyword query k_q




Fig. 4. Node-partition graphs.

"sea food" supplied by the user. In the first iteration, node k_2 retains 0.5 amount of ink and distributes 0.5 amount of ink to its neighbor document nodes d_1-d_3 according to the adjusted edge weights. In the second iteration, d_3 distributes its active ink of amount 0.325 to its neighbor keyword query nodes k_2 and k_3 . BA terminates at the sixth iteration where the active ink of each node is smaller than ϵ . The top-1 suggestion (excluding user query k_2) is $k_3 = "lobster"$, with the largest amount of retained ink (0.098).

3.2 Partition-Based Algorithm

Algorithm BA can be slow for several reasons. First, at each iteration, only one node is processed; thus, the active ink drops slowly and the termination conditions are met after too many iterations. Second, given the large number of iterations, the overhead of maintaining queue Q is significant. Finally, the nodes distribute their active ink to all their neighbors, even if some of them only receive a small amount of ink. To improve the performance of BA, in this section, we propose a partition-based algorithm that divides the keyword queries and the documents in the KD-graph G into groups. Let $\mathcal{P}^K = \{P_i^K\}$ be the partitions of the keyword queries and $\mathcal{P}^D = \{P_i^D\}$ be the document partitions. Algorithm PA follows the basic routine of algorithm BA, but with the following differences:

- 1) *Node-partition graphs.* PA uses two directed graphs G^{KP} and G^{DP} constructed offline from the KD-graph G and partitions \mathcal{P}^K and \mathcal{P}^D . In graph G^{KP} , a keyword query node k_i connects to a document partition P^D if k_i connects in G to at least one document in P^D . Similarly, in graph G^{DP} , a document node d_j connects to a keyword partition P^K if d_j connects in G to at least one keyword query node k_i . As an example, in Fig. 4, the document partitions are $P_1^D = \{d_1, d_2\}$ and $P_2^D = \{d_3, d_4, d_5\}$ and the keyword query partitions are $P_1^K = \{k_1\}$ and $P_2^K = \{k_2, k_3\}$. The edge weights are defined based on graph G_q , computed during the execution of PA. Each edge weight shown in Fig. 4 indicates the portion of the ink to be distributed to a partition P from a node v that is the sum of the adjusted weights of the edges from node v to the nodes in P according to G_q .




Fig. 5. Ink distribution methods.

- 2) *Ink distribution.* In PA, each node distributes its active ink to its neighbor partitions (contrast this to BA, where each node distributes its active ink to each of its neighbor nodes). The priority queue used in BA maintains the nodes that will distribute ink, but the priority queue used in PA records the partitions that will be processed. The ink received by a partition is not spread to the nodes inside the partition until this partition reaches the head of the priority queue. The benefit is that a partition may receive ink from the same node several times while waiting in the queue, so that the nodes in this partition receive ink in batch when this partition reaches the head of the queue. In algorithm PA, the active ink drops fast and the termination conditions may be fulfilled early. Thus, the number of iterations needed is largely reduced and so is the cost spent for maintaining the priority queue Q . Moreover, since the number of partitions is much smaller than that of nodes, the size of queue Q is much smaller compared to that used in BA, so operations on it are fast as well. As an example, in Fig. 5, in algorithm BA, node k_2 distributes its active ink to each of its three neighbor nodes d_1-d_3 . However, in algorithm PA, the active ink of k_2 is only distributed to two recipients: partitions P_1^D and P_2^D ; an underlying document node will not receive the ink, until its partition reaches the top of the queue.
- 3) *Lazy distribution mechanism.* In BA, a node distributes ink aggressively, i.e., each of its neighbor nodes receives ink no matter how much it is. On the other hand, in algorithm PA, we adopt a lazy distribution mechanism that relies on threshold ϵ . If the amount of the ink to be distributed from a node v to a partition P is smaller than ϵ , P does not receive the ink immediately; instead, the ink is accumulated (i.e., buffered) at v . Later, if at some point the ink accumulated at v for partition P exceeds ϵ , P receives it. Overall, this lazy distribution mechanism delays the distribution of small amounts of ink across the graph that would otherwise result in many updates, reducing the computational cost significantly. As a toy example in Fig. 5b, the amount of ink (0.07) to be distributed from node k_2 to partition P_1^D waits at k_2 when $\epsilon = 0.1$.

Algorithm 2 is a pseudocode for PA. Priority queue Q maintains the partitions to be processed in descending order of their keys (line 1). A partition waiting in the queue may have received ink from multiple nodes. Let $P.\text{aink}_{v_i}$ be the ink received by partition P from node v_i . The key of P in the queue is the maximum value of $P.\text{aink}_{v_i}$. Priority queue C stores the candidate suggestions in descending order of their retained ink, initialized as empty (line 1). Q initially contains one entry referring to the keyword query




Fig. 6. Illustration of Algorithm PA.

partition P containing the user supplied keywords with active ink 1 (line 2). The sum of the active ink of all nodes $AINK$ is set to 1 (line 3). The termination conditions are the same as those of algorithm BA (lines 4 and 8). For a partition P_t being processed, PA firstly spreads its received ink to P_t 's nodes (line 10). When processing a keyword query partition, each keyword query node v inside retains α portion of its active ink (line 15). The ranking of this keyword query node in C is updated and the active ink $AINK$ is modified (line 16). For each document partition P_i connected from v in G^{KP} , if the amount of accumulated ink for P_i at v exceeds ϵ , partition P_i receives the ink. Otherwise, the ink is accumulated at v for P_i (lines 24-32). Similarly, when processing a document partition, for each document node v inside the partition, the lazy ink distribution mechanism is applied. As before, the only difference is that document nodes do not retain any ink. The algorithm finally returns the top- m candidate suggestions in C as the result (line 33).

Note that the partitioning is done offline, and thus it does not add any cost to the online query suggestion process. In PA, the cost of maintaining the partition information (in terms of both the memory usage and the CPU cost) is negligible when the number of partitions is small (e.g., 16-64 as shown in our experiments).

Example 3. We illustrate algorithm PA (for $m = 1$, $\epsilon = 0.1$, and $\alpha = 0.5$) using our running example ($k_q = "seafood"$ and $\lambda_q = (0.2, 0.2)$). Fig. 4 shows the node-partition graphs. Fig. 6 shows how the scores (ink) of the keyword queries are computed. Initially, partition P_2^K contains the user supplied keyword k_2 having active ink 1. In the first iteration, partition P_2^K is processed. Each keyword query node inside partition P_2^K retains α portion of its active ink. In this example, k_2 retains 0.5 amount of ink. Document partitions P_1^D and P_2^D connected from k_2 according to graph G^{KP} receive 0.175 and 0.325 amount of ink, respectively. In the second iteration, document partition P_2^D is processed. Each node inside P_2^D first receives the ink from k_2 and then distributes ink to the connected keyword query partition P_2^K according to graph G^{DP} . In the third iteration, partition P_2^K is processed. Each node inside P_2^K receives the ink from d_3 . After that, node k_3 accumulates ink for partition P_2^D , since the amount of ink

for P_2^D is 0.097 ($< \epsilon$). Similarly, node k_2 accumulates ink for partitions P_1^D and P_2^D . In the fourth iteration, partition P_1^D is processed. For each node inside P_1^D , the ink to be distributed does not exceed threshold ϵ , so that d_1 and d_2 accumulate ink for partitions P_1^K and P_2^K . In the end, the top-1 suggestion is keyword query $k_3 = "lobster"$.

Algorithm 2. PA

```

Input:  $G(D, K, E), G^{KP}, G^{DP}, q = (k_q, \lambda_q), m, \epsilon$ 
Output:  $C$ 
1 PriorityQueue  $Q \leftarrow \emptyset, C \leftarrow \emptyset$ ;
2 Add partition  $P \ni k_q$  to  $Q$  with  $P.aink \leftarrow 1$ ;
3  $AINK \leftarrow 1$ ;
4 while  $Q \neq \emptyset$  and  $Q.top.aink_{v_i} \geq \epsilon$  do
5   Deheap the top entry  $P_t$  from  $Q$ ;
6    $tm =$  the top- $m$  entry from  $C$ ;
7    $tm' =$  the top- $(m + 1)$  entry from  $C$ ;
8   if  $tm.rink > tm'.rink + AINK$  then
9     break;
10  Spread the active ink to nodes in  $P_t$ ;
11  for each node  $v$  in partition  $P_t$  do
12     $distratio = 1$  ;
13    if  $v$  is a keyword query node then
14       $distratio = 1 - \alpha$  ;
15       $v.rink \leftarrow v.rink + v.aink \times \alpha$ ;
16       $AINK \leftarrow AINK - v.aink \times \alpha$ ;
17      if there exist a copy  $t$  of  $v$  in  $C$  then
18        Remove  $t$  from  $C$ ;
19         $v.rink \leftarrow v.rink + t.rink$ ;
20      Add  $v$  to  $C$ ;
21      Get partition set  $\mathcal{P}$  connected from  $v$  in  $G^{KP}$ ;
22    else
23      Get partition set  $\mathcal{P}$  connected from  $v$  in  $G^{DP}$ ;
24      for each partition  $P_i$  in  $\mathcal{P}$  do
25         $ink \leftarrow v.aink \times distratio \times \tilde{w}(v, P_i)$ ;
26        if  $ink + v.acc.P_i \geq \epsilon$  then
27           $P_i.aink \leftarrow ink + v.acc.P_i$ ;
28          if there exist a copy  $P'_i$  of  $P_i$  in  $Q$  then
29            Remove  $P'_i$  from  $Q$ ;
30             $P_i.aink \leftarrow P_i.aink + P'_i.aink$ ;
31          Add  $P_i$  to  $Q$ ;
32        else
33          Accumulate ink at node  $v$  for  $P_i$  ( $v.acc.P_i$ );
33 return the top- $m$  entries (excluding  $k_q$ ) in  $C$ ;

```

Partitioning methods. We consider four partitioning methods for keyword query/document nodes, to be empirically evaluated in Section 4.

Random partitioning. Keyword queries are evenly and randomly partitioned into a predetermined number of partitions. The same is done for the documents.

Spatial partitioning. First, a regular grid is used to partition the euclidean space of document locations. The documents whose locations lie in the same grid cell form a partition. The keyword queries are partitioned according to the document partitions. The idea is inspired by the *duality of word and document clustering* [29]. Specifically, let $\mathcal{P}^D = \{P_1^D, P_2^D, \dots, P_n^D\}$ be the document partitions. We initialize N empty keyword query partitions $\mathcal{P}^K = \{P_1^K, P_2^K, \dots, P_N^K\}$. According to graph G^{KP} , which connects keywords to document partitions, each keyword query node k_i is connected to a set of document partitions $P(k_i) = \{P_j^D\}$. Let P_j^D be the document partition connected from k_i with the highest edge weight, i.e., $\arg \max_{P_j^D \in P(k_i)} w(k_i, P_j^D)$. Keyword query node k_i is added to partition P_j^K that has the same subscript j as partition P_j^D . In the end, some of the keyword query partitions in \mathcal{P}^K might be empty, and thus are removed.

Textual partitioning. Each document d is associated to a vector $d.\psi$ where each dimension refers to a keyword query k connected to d in the KD-graph G and the value of this dimension is the edge weight $w(d, k)$. Let $\cos(d_i.\psi, d_j.\psi)$ be the cosine similarity between the vectors of documents d_i and d_j . A clustering algorithm, e.g., k-means, is applied on the document vectors using the cosine similarity, so that the document partitions are obtained. The keyword queries are partitioned based on the document partitions in the same way as in the spatial partitioning method.

Hybrid partitioning. Let $f_H(d_i, d_j) = \gamma \times \cos(d_i.\psi, d_j.\psi) + (1 - \gamma) \times \text{dist}(d_i.\lambda, d_j.\lambda)$ be a hybrid distance between documents d_i and d_j , considering both their euclidean distance and cosine similarity. A clustering algorithm, e.g., k-means, is applied on the documents using the hybrid distance, so that the document partitions are obtained. The keyword queries are partitioned as in the previous methods.

4 EMPIRICAL STUDY

We conducted extensive experiments to evaluate the effectiveness and efficiency of LKS. Section 4.1 presents our experimental setup. The effectiveness of our LKS framework compared to query suggestion that does not consider locations is evaluated in Section 4.2. The runtime performance of algorithms BA and PA is evaluated in Section 4.3.

4.1 Setup

Two datasets AOL and TWEET are used in our evaluation. AOL is based on a real query log of the AOL search engine [30]. Each record in the log contains a keyword query, the time when the query was submitted, and a URL clicked. Since the locations of the URLs are not readily available in the query log, we geo-located the URL domains with the help of the Freegeoip⁶ project and removed the URLs

TABLE 1
Parameters

Description	Parameter	Values	Default Value
Number of partitions	N	1,4,16,64,256	16
RWR probability	α	0.2,0.35,0.5, 0.65,0.8	0.5
Edge weight adjustment param.	β	0,0.25,0.5, 0.75,1	0.5
Threshold ($\times 10^{-5}$)	ϵ	0.1,0.5,1, 5,10	1
Number of documents (M)	$ D $	0.5, 1, 1.5, 2, 2.5	1.5

without any geo information. In addition, we cleaned the query log by removing the keyword queries without click information and with frequency less than 3. We added an edge between a keyword query node k_i and a URL node d_j if there exist records containing both k_i and d_j in the query log. The edge weight was defined by the number of records containing both k_i and d_j (i.e., the click count). At last, the constructed KD-graph G has 629,875 keyword query nodes, 496,221 document nodes (URLs), and 2,778,050 edges. Dataset TWEET is based on 3,198,266 tweets published inside the New York area, collected using Twitter's Streaming API.⁷ Each tweet has a text message and a location where the tweet is posted. Following the methodology in [31], we extracted phrases of length 1 to 10 from the text messages and used them to model the keyword queries. Only the phrases ending with either a noun or an adjective and with frequency at least 3 are kept, in order to reduce the number of noisy queries. We add an edge between a keyword query node k_i and a tweet node d_j if d_j contains keywords k_i . The edge weight is the tf-idf score of k_i w.r.t. d_j . Afterwards, documents with no connected keyword query nodes are removed. Finally, the constructed KD-graph has 781,465 keyword queries, 1,482,064 documents (tweets), and 12,078,958 edges.

From each dataset, we randomly selected 100 keyword queries and used them as a workload. For each query, the user location is considered to be the location of a randomly selected document from those that contain the query keywords. The reasoning behind this is that the query location distribution typically tends to follow the distribution of the documents that contain the query.

All tested methods were implemented using Java. The experiments were run on a machine with Intel Core i7-3770 3.40 GHz and 16 GB main memory. Our LKS framework and the algorithms are evaluated under various parameter settings, as shown in Table 1. For each parameter, we used a wide range of values in order to test their effects. The number of suggestions m is fixed to 5 (the performance of the evaluated algorithms is not sensitive to it). Finally, the default partitioning method for PA is Spatial Partitioning.

4.2 Effectiveness of the LKS Framework

In this section, we assess the effectiveness of LKS using real examples, a user study, and quantitative measurements. These evaluation metrics verify whether the keyword

6. <http://freegeoip.net>

7. <https://dev.twitter.com/docs/api/streaming>

Authorized licensed use limited to: SHENZHEN UNIVERSITY. Downloaded on January 28, 2026 at 00:53:37 UTC from IEEE Xplore. Restrictions apply.




Fig. 7. User Supplied Keyword $k_q = \text{"airport"}$, Query Location λ_q near LaGuardia Airport.

suggestions are (i) semantically relevant to the original query and (ii) able to find documents that are close to the query location.

4.2.1 Case Study

We first demonstrate the effectiveness of considering the user location in query suggestion, by showcasing query examples on the TWEET data set. Consider a user-supplied keyword query with $k_q = \text{"airport"}$. Figs. 7 and 8 show the top-1 suggestion for k_q considering two different query locations λ_q and using various values of β . In all plots, the triangle denotes the query location, the crosses denote the *processed documents* by the PA algorithm (i.e., the documents that distribute ink during the course of the algorithm), and the circles are the retrieved document by the suggested keywords close to the query location (located within 10km from the query location). When $\beta = 1$, LKS ignores the query location; thus, no matter where the query location is, the suggestion for the user supplied keyword "airport" is always "LaGuardia Airport" (the most relevant phrase to "airport" when disregarding location data) as shown in Figs. 7a and 8a. When the query location is close to LaGuardia Airport, the suggested keywords "LaGuardia Airport" can retrieve some relevant documents nearby (the circles in Fig. 7a). However, when the query location is close to

Newark, there are no relevant documents around related to the suggestion (no circle in Fig. 8a).

When $0 < \beta < 1$ ($\beta = 0.5$ in this example), the suggestion depends on the query location. Keywords "LaGuardia Airport" are suggested when the query is issued near LaGuardia (Fig. 7b), while "Newark Liberty International Airport" is the top suggestion when the query location is close to Newark (Fig. 8b). Each of the two suggestions can retrieve relevant documents (circles) close to the corresponding query location. When $\beta = 0$, the suggestion only depends on the query location and can fetch relevant documents nearby as shown in Figs. 7c and 8c. However, sometimes the suggestion is less specific (i.e., less relevant) compared to the case of $\beta = 0.5$, e.g., "Newark Liberty International Airport" versus "Newark Liberty". This is because the spatial-only recommendation favors the queries having more relevant documents close to the query location, regardless of their textual relevance. When the query location is near JFK, "John F Kennedy International Airport" and "Kennedy International" are reported when β is set to 0.5 and 1.0, respectively. We omit the details due to space constraints.

We also observe that for small β values, fewer documents are involved in the computation (crosses) because more weights are given to the nearby documents. For example, the number of crosses is 2,296, 2,177, and 822 in Figs. 8a and 8b, and c, respectively. Moreover, small β results in more




Fig. 8. User Supplied Keyword $k_q = \text{"airport"}$, Query Location λ_q near Newark Liberty International Airport.

Authorized licensed use limited to: SHENZHEN UNIVERSITY. Downloaded on January 28, 2026 at 00:53:37 UTC from IEEE Xplore. Restrictions apply.




Fig. 9. Suggestions to query $k_q = \text{"prada"}$, $\lambda_q = (0.75, 0.7)$.

general query suggestions, which retrieve more nearby documents. As an example, the number of circles is 600 and 630 in Figs. 8b and c, respectively.

A case where LKS is especially useful is when there are no (or very few) relevant documents near the user location. LKS recommends to the user alternative query keywords, which match the user's intention and at the same time find nearby documents. For example, Fig. 9 present the top recommendation to $k_q = \text{"prada"}$ on our AOL dataset for various β values. When the location information is ignored ($\beta = 1$, Fig. 9a), query "prada store" is recommended, which however retrieves no nearby document. LKS with $\beta = 0.5$ (Fig. 9b) recommends a different brand "gucci", which can retrieve three related documents. Finally, when the model is over-biased to the location factor ($\beta = 0$, Fig. 9c), the suggestion "handbags" retrieves even more (8) nearby documents, which however is less irrelevant to k_q .

4.2.2 User Study

As an additional way to test the effectiveness of LKS, we randomly selected 100 AOL queries with $\beta = 0$ (suggestion based on location only), $\beta = 0.5$ (our LKS model) and $\beta = 1$ (previous work, i.e., location-independent models), and distributed them to seven colleagues. To ensure the fairness of the user study, the participants were not informed about the details of this project and the respective β setup of the three scenarios. For each input query, we showed the participants (1) the top suggestion and (2) the top-5 documents retrieved by spatial keyword search using the *boolean range query* model [32], so that they are within 0.1 distance to the user location and are ranked by their relevance to the suggestion. Then, we asked the participants to suggest two scores, (1) S_K : how relevant (in 0-5) the suggestion is to the original input keyword, and (2) S_D : for each document, judge if it is related to the initial query intent (0 or 1 for each document, in case fewer than five documents are within 0.1 distance, set the remaining slots to 0, which represents a penalty if the suggestion still cannot retrieve at least five nearby results). S_D is the sum of these five binary values. S_K verifies textual relevance of the suggestion. On the other hand, S_D verifies usefulness of the suggestion by the relevance of the retrieved nearby documents (i.e., verifies the user's intention both semantically and spatially). The average S_K

and S_D scores are presented in Fig. 10, showing the superiority of LKS. Note that document relevance S_D is low at $\beta = 1$ (previous work) because without location information, the suggested keywords may fail to find five nearby documents; S_D is also low at $\beta = 0$ because although many documents are retrieved, they are not very relevant to the initial input.

4.2.3 Quantitative Study

Designing quantitative evaluation metrics for LKS is challenging, because it is hard to establish the ground truth for suggestions. Considering the two criteria of good suggestions, we evaluate (i) the semantic relevance of the suggested keyword queries (together with their retrieved nearby documents) w.r.t. the user's initial query and (ii) the number of nearby documents retrieved by the query suggestions. As a competitor to our LKS framework, we implemented the *influence tag co-occurrence* (INF) method proposed in [19]. INF is designed to employ both spatial and textual information into tag recommendation for Flickr and has been shown to be more effective than alternative approaches [19]. Given a photo p published at $p.\text{loc}$ and a tag t , the task is to retrieve k tag recommendations to t . The idea is that if a tag t' co-occurs with t in a photo p' , they are considered textually relevant. In the meantime, if p' is close to p (i.e., $\text{dist}(p'.\text{loc}, p.\text{loc})$ is small), the relevance between t' and t is even higher. In INF, the relevance between t and t' is defined according to a influence score

$$\text{inf}(t, t') = \frac{\sum_{p' \in P \cap P'} 2^{-\frac{\text{dist}(p', p)}{r}}}{|P \cup P'|}. \quad \text{Here } P \text{ and } P' \text{ are the sets of photos containing tag } t \text{ and } t', \text{ respectively, and } r \text{ is a}$$




Fig. 10. User study on AOL.




Fig. 11. Number of retrieved nearby documents.

parameter to adjust the importance of spatial distance.⁸ To apply INF in keyword suggestion, we consider the co-occurrence of keywords in the same document and exploit a similar relevance function.

We first report the average number of nearby documents retrieved by (a) the original query, (b) the queries suggested by INF, and (c) the queries suggested by our LKS framework, over the workload of TWEET when varying parameters α (in Equation (3)) and β (in Equation (1)). Fig. 11a plots the number of documents retrieved by the original input and the keyword queries suggested by INF or LKS, within a ratio ρ of the maximum euclidean distance between any pair of documents ($\rho = \{0.05, 0.1\}$). The queries suggested by INF can retrieve some more nearby locations. However, the number of documents retrieved by the LKS-suggested queries is significantly higher than that of either the original input, or the INF suggested keyword queries, for all α values, showing the value of LKS w.r.t. the location criterion. Obviously, for larger ρ values more documents are retrieved. Also, the number of documents decreases as α increases. The reason is that for small α values, the distributed ink reaches more documents compared to the case where α is large. Hence, more document nodes near the query location participate in the ink distribution process and more relevant queries are reached and included in the top- m ranking. As special cases, when $\alpha = 1$, no keyword queries can be returned because all the ink is retained by the original query, while when $\alpha = 0$, no restart is involved in the random walk, hence the result is independent of the input. Fig. 11b reports the number of documents retrieved by the original input and the keyword queries suggested by INF or LKS when varying β . The number of documents retrieved by the queries suggested by LKS decreases as β increases. This is because a large β weighs the user location low; the textual relevance criterion dominates the spatial closeness criterion, therefore fewer nearby documents are relevant to the suggestions. Note that our experiments on AOL show similar results and are thus omitted.

Next, we evaluate how relevant the suggested keyword queries are to the user's information needs. To begin with, we first evaluate the textual relevance between the original input and the keyword queries suggested by INF or LKS,

8. We tested various r values from 0.1 to 0.0001 and observed that with decreasing r , more nearby documents are retrieved, but the semantic relevance of the suggested keywords declines. We present only the results with $r = 0.001$, which retains the best balance between the two factors.

Authorized licensed use limited to: SHENZHEN UNIVERSITY. Downloaded on January 28, 2026 at 00:53:37 UTC from IEEE Xplore. Restrictions apply.




Fig. 12. Textual relevance evaluation.

following a similar approach used in [33], [34], [35]. Specifically, we utilize the Open Directory Project database (ODP, a.k.a. DMOZ),⁹ which provides a comprehensive human edited directory of the web. ODP can match a given query to a hierarchy of categories. For example, the query "Hard Disk" matches category "Computers: Hardware: Storage: Hard Disk" and the query "Memory" matches "Computers: Hardware: Components: Memory". Hence, the relevance between two queries can be evaluated by the similarity between their corresponding categories C_1 and C_2 in the hierarchy. Similarity is measured by the length of the longest common path prefix $P(C_1, C_2)$ normalized by the length of the longer path in the hierarchy, i.e., $sim(C_1, C_2) = \frac{|P(C_1, C_2)|}{\max\{|C_1|, |C_2|\}}$. For example, the similarity between "Hard Disk" and "Memory" is 2/4 as the longest common path "Computers: Hardware" has length 2. In our experiments, we evaluated the relevance between the original input and the top suggested keyword queries by INF or LKS, based on their most similar categories provided by ODP. The results are presented in Fig. 12a. The first observation is that the LKS-based suggestions are generally more promising than INF based ones, unless β is close to 0 (spatial only). This is because INF relies merely on co-occurrence to judge textual relevance, which cannot capture many hidden semantic relationships between keyword queries. LKS, on the other hand, captures textual proximity based on the whole structure of KD-graph and thus retrieves more reliable suggestions. In addition, we also observe that LKS achieves best ODP similarity when β is selected between 0.25 and 0.75.

After the direct evaluation of suggested query keywords in the previous experiment, we now evaluate the nearby documents retrieved by them. Let $\{d_i^o\}$ be the set of top-10 documents retrieved by the original query and $\{d_i^s\}$ be the set of top-10 documents retrieved by a suggested query within a spatial range ($\rho = 0.1$), where i indicates the ranking position of a document. We compute the cosine similarity $\cos(d_i^o, d_i^s)$ between the documents with the same ranking position from the two sets $\{d_i^o\}$ and $\{d_i^s\}$. We adopt the nDCG measure of [36] to aggregate these cosine similarities so that the document similarity is reduced logarithmically, proportionally to its ranking position

$$COS = \frac{\cos(d_1^o, d_1^s) + \sum_{i=2}^{10} \cos(d_i^o, d_i^s) / \log(i)}{1 + \sum_{i=2}^{10} 1 / \log(i)}.$$

9. <http://www.dmoz.org>




Fig. 13. Response time when varying ϵ ($\times 10^{-5}$).

The higher the value of COS is, the closer the nearby documents are to user's information needs. Besides, we also compare the textual relevance (i.e., $tf \times idf$) of documents $\{d_i^s\}$ and $\{d_i^o\}$ w.r.t. the original query. Let s_i^o and s_i^s be the textual score of document d_i^o and d_i^s w.r.t. the original query, respectively. We take the ratio of s_i^s to s_i^o for the same ranking position. Similar to COS , these ratios are aggregated as TS . Fig. 12b shows the average COS and TS over the workload of TWEET when varying β . The results are mostly consistent with Fig. 12a. LKS again has better textual quality than INF unless β is very small (< 0.25). In the meantime, LKS achieves best effectiveness at non-extreme β values, where it takes advantage of both the textual relationships from the original KD-graph, and the local intent of the input query.

Considering the above factors, we conclude that an average value of β (e.g., $0.25 < \beta < 0.75$) should be used in order for the keyword query suggestions to have a sufficient number of documents near the user location as results, which also meet the information needs of the user. In summary, our LKS framework, which balances between relevant and nearby search results, is very effective, based on our experimental analysis.

4.3 Efficiency

In this section, we evaluate the performance of BA and PA on the two datasets under various settings and report their average runtime over the workload of 100 queries.¹⁰

4.3.1 Varying ϵ

Both BA and PA follow the general idea of BCA [25], and their runtime depends on parameter ϵ . The smaller ϵ is, the smaller the error incurred by BCA (and BA/PA) on computing the PageRank scores on G_q . We find that when $\epsilon \leq 10^{-8}$, the top query suggestions become stable. Thus, we consider the top suggestions calculated using $\epsilon = 10^{-8}$ as the base for computing the approximation error of the top suggestions calculated using $\epsilon > 10^{-8}$. Following [28], the error is computed as $1 - AP$ where AP is the average precision of the retrieved results compared to the true ones. An appropriate ϵ value should be selected so that (1) the computation can be finished in real time and (2) the error rate is reasonably low. Fig. 13 shows that the response time of BA drops dramatically as ϵ increases, since a large ϵ significantly reduces the number of iterations in BA and saves many computations. However, the approximation error is high for large values of ϵ . On the other hand, PA runs fast even for small




Fig. 14. Response time when varying α .

values of ϵ , for which the approximation error is low. PA outperforms BA by 1-2 orders of magnitude when ϵ is smaller than 10^{-6} . Therefore, under the same time constraint (e.g., < 1 s), PA can support smaller ϵ values, and can thus guarantee a lower error rate.

Because of ϵ , BA and PA may produce different suggestions, but the differences are mainly in the low-ranked suggestions. In practice, users only consider the highly ranked suggestions. In our experiments, the top-5 suggestions offered by BA and PA are identical in 98 and 99 percent of the times on AOL and TWEET, respectively.


4.3.2 Varying α

Parameter α in Eq. (3) indicates the random walk restart probability that corresponds to the proportion of ink retained on a query node. As shown in Fig. 14 the response times of both BA and PA decrease as α increases. This is because for larger α , more amount of ink is retained at the keyword query nodes; thus, the amount of active ink drops quickly and the termination condition is satisfied early. We observe that PA is more robust to α and outperforms BA significantly when α is small. In the special case, when $\alpha = 1$, all the ink will be retained on the query node in the beginning, thus both algorithms terminate immediately returning no results. In the other extreme case, when $\alpha = 0$, no ink is retained in any node at each step, therefore the ink keeps being redistributed until the random walk process converges to a stable state. In this case, the final scores of the nodes only depend on the structure of the graph, and not on the starting node (query node); no matter what query is given, the suggestions are always the same (i.e., similar to global PageRank scores). Therefore, both extreme cases ($\alpha = 1$ or 0) cannot give useful results and are not considered.

4.3.3 Varying β

Parameter β in Fig. 15 shows the response time when varying β . Both BA and PA are bell-shaped. When β approaches 0 or 1, only the documents that are either close to the query location or have high textual relevance to the user supplied keywords are involved in the computation. However, when β is close to 0.5, the response time increases slightly, since more documents are considered, i.e., the ones that are a bit far from the query location but with high textual relevance and the ones that are close to the query location but with slightly low textual relevance. Algorithm PA outperforms BA for all values of β by a wide margin.

10. Experiments with larger workloads gave us similar results.

Fig. 15. Response time when varying β .


4.3.4 Varying Partitioning Methods and Number of Partitions

We evaluate the four partitioning methods introduced in Section 3.2 using various number of partitions on the two datasets. The results are shown in Figs. 16a and 16b. The performance of PA seems to be insensitive to which partitioning approach is used. This can be attributed to the fact that the KD-graph and its corresponding matrices are very sparse. Therefore, the difference between the clusters formed by these approaches is not significant enough to influence the performance of PA. To verify this assertion, we conducted experiments using two denser versions of our datasets; the dense AOL_D is built by selecting the top-50 percent documents with the most connecting edges from AOL, and the dense TWEET_D is built by selecting the top-30 percent documents from TWEET. To make the graphs even denser, we employ *item-based collaborative filtering* (IBCF) [37] to estimate additional edge weights from queries to documents (queries correspond to “users” and documents correspond to “items”). We first calculate the similarity among nearby documents (within 1 percent distance threshold), based on their co-connected queries and weights. Afterwards, for each non-existent keyword to document edge (k, d) , we estimate its weight by considering documents D_s that are similar to d and the weights of edges connecting k to D_s . Formally,

$$\omega_e(k, d) = \frac{\sum_{d' \in D_s} sim(d, d') \times \omega(k, d')}{\sum_{d' \in D_s} sim(d, d')}.$$
 (4)

If the estimated weight $\omega_e(k, d)$ is again zero, we keep the pair (k, d) unconnected. After the above procedure, we have AOL_D containing over 26 M edges among 328 K queries and 325 K documents, and TWEET_D containing over 38 M edges among 302 K queries and 477 K documents. Compared to our original datasets, AOL_D is 27 times denser than AOL and TWEET_D is 17 times denser than TWEET. The experimental results of Figs. 16c and 16d confirm that with denser matrices, the partitioning methods do have an effect on performance. Specifically, the hybrid method outperforms all other approaches because it takes advantage of both the spatial and textual factors during the ink propagation procedure, and thus predicts better how the ink may tend to flow and cluster, achieving better partitioning.

Note that the number of partitions N , directly affects the performance PA performs best when the number of partitions is 64. When N is small, the lazy distribution mechanism has little effect, since the size of each partition is large

Fig. 16. Varying partitioning methods and N .

and there is little chance that the ink is being accumulated. Thus, the cost saved by the ink accumulation is limited. On the other hand, when N is too large, the lazy distribution mechanism is applied for too many partitions and the bookkeeping overhead at each node for the accumulated ink per partition is high, negatively affecting the performance of PA and counterbalancing the savings due to the lazy distribution.

4.3.5 Varying Data Scale

To test the effect of data size, we constructed a large TWEET_US dataset from about 10M tweets published in general US. area, in similar way to the TWEET dataset. We use subsets of TWEET_US with various numbers of documents / keywords. Fig. 17 reports the response time of BA and PA when $|D|$ varies from 0.5 to 2.5 M. Note that with different $|D|$, the corresponding keyword node set size $|K|$ also changes, as shown on the top of Fig. 17. Based on the experiments, PA constantly outperforms BA by three-four times with various $|D|$ or $|K|$ settings.

5 RELATED WORK

Related work on query suggestion is discussed in Section 5.1. Techniques for RWR computation are reviewed in Section 5.2.




Fig. 17. Varying number of documents and keywords.

5.1 Keyword Query Suggestion

Keyword query suggestion approaches can be classified into three main categories: random walk based approaches, learning to rank approaches, and clustering based approaches. We also briefly review alternative methods that do not belong to any of these categories. To the best of our knowledge, no previous work considers user location in query suggestion.

5.1.1 Random Walk Based Approaches

The methods in this category use a graph structure to model the information provided by query logs, and then apply a random walk process on the graph to compute the suggestions. Craswell and Szummer [4] apply such an approach on the query-click graph and suggest queries based on personalized PageRank scores. Mei et al. [5] rank queries based on hitting time, which reflects the probability that a random walker arrives a node within certain steps. Song and He [6] combine both the clicked and skipped URLs from users in the query-URL bipartite graphs in order to also consider rare query suggestions (using clicked URLs only favors popular queries). This work was extended in [38] to achieve diversification for the recommended queries; the suggested queries are re-ranked in a way that maximizes the diversification function. Zhu et al. [39] generate query recommendations by considering diversity and relevance in a unified way during RWR search. Miyanishi and Sakai [7] introduce Time-aware Structured Query Suggestion (TaSQS), an algorithm for presenting query suggestions along a timeline. Our proposed LKS framework is orthogonal to and can be easily integrated in all the suggestion methods that use the query-URL bipartite graph.

Boldi et al. [9] proposed an approach that applies RWR from the original query or from a small set of queries representing the recent querying history of the current user on a *query-flow* graph. Deng et al. [35] investigate and develop an entropy-biased framework for modeling click graphs, in which common clicks on frequent and general pages are given less weight than those on less frequent but more specific ones. Anagnostopoulos et al. [8] formulate the query-recommendation problem as a decision problem on how to perturb the transition probabilities between queries in the query-flow graph so as to maximize the expected utility of a random walk. Song et al. [10] mine a term-transition graph from search engine logs and apply a topic-based unsupervised Pagerank model that suggests queries based on the topic distribution and term-transition probability within each topic. The idea of our LKS framework can be applied to methods that use a query-flow graph, since the query history may also have location information. This is an interesting subject for our future work.

5.1.2 Learning to Rank Approaches

Some query suggestion approaches [40] are based on learning models trained from co-occurrences of queries in search logs. Another learning-to-rank approach [41] is trained based on several types of query features, including query performance prediction. Li et al. [11] train a hidden topic model. For each candidate query, its posterior distribution over the hidden topic space is determined. Given a user query q , a list of suggestions is produced based on their

similarity to q in the topic distribution space. Our work is not based on learning models; in the future, it would be interesting to study how these models can be extended to consider location information.

5.1.3 Clustering Based Approaches

Beferman and Berger [2] view the query log as a query-URL bipartite graph. By applying an agglomerative clustering algorithm on the vertices in the graph, query clusters can be identified. Then, given a user-supplied query q , the queries that belong to the same cluster as q does are returned to the user as suggestions. Wen et al. [42] further extended the approach to also take into account the similarity between the query content during clustering. In [1], a similar approach is proposed: the queries are modeled as term-weighted vectors and then clustered. The vector of a query q includes the clicked URLs by the users who posed q as terms and the weights are calculated based on term frequency and the click popularity of the URL in the answers of q . Cao et al. [3] take into account the immediately preceding queries as context in query suggestion. They summarize queries in search logs into concepts by clustering a query-URL bipartite graph. User session data are converted to concept sequences and indexed by a suffix tree. The query sequence submitted by the user is mapped to a sequence of concepts; the suffix tree is then searched to find query suggestions. Finally, Li et al. [43] cluster queries from search logs to extract query concepts, based on which recommended queries are selected and employ a probabilistic model and a greedy heuristic algorithm to achieve recommendation diversification. Location information could also be considered in all these clustering models. Such an approach is out of the scope of our current work, but we are interested in investigating its effectiveness in the future.

5.1.4 Miscellaneous Approaches

Zhang and Nasraoui [44] create a graph with edges between consecutive queries in each session, weighted by the textual similarity between these queries. A candidate suggestion for a given query is given a score based on the length of the path between the two queries, aggregated across all sessions in a query log where the query and the suggestion co-occurred. Cucerzan and White [45] generate query suggestions based on user *landing pages* (that is, the webpages that users end a query with, through post-query browsing). Given a user query, they utilize its recorded landing pages and suggest other queries that have these landing pages in their top ranked results. A probabilistic mechanism [31] generates query suggestions from the corpus without using query logs.

5.1.5 Location-Aware Type-Ahead Search (LTAS)

References [17] and [18] both study the problem of location-aware type-ahead search, also known as *instant search*. LTAS finds documents near a user location, as the user types in a keyword query character by character. This problem is more related to keyword query completion than to the query suggestion problem that we study in this paper, since the recommended keywords must have the user's

input as prefix. On the other hand, the query suggestion problem that we study in this paper takes a completed query and recommends other queries that are semantically relevant without the constraint that the suggestions should have the original user query as prefix. Therefore, our LKS framework is more flexible and can help users to express various aspects of a topic. The suggested keywords can be different than the user-supplied keywords, but they should be textually relevant. In addition, the methods for LTAS are very different to our LKS algorithms, as they take advantage of the prefix requirement to reduce the search space (with the help of trie data structures).

5.1.6 Location-Aware Suggestions Based on User History

Google [46] provides location-based query suggestions by simply selecting the user's past search queries that have results close to the user's current location. These suggestions may be insufficient if the user did not perform any historical searches near her current location. In addition, query suggestion based on location only may not match the user's search intent. On the other hand, our framework aims at suggesting keyword queries that satisfy the user's information needs and have nearby results, irrespectively to the user's search history.

5.1.7 Query Relaxation

A relevant problem to query suggestion in relational databases is called *query relaxation*. The objective is to generalize an SQL query in case of too few or no results [47]. Query relaxation approaches cannot be applied for keyword query suggestion, because they require the relaxed query to contain the results of the original query, which is not essential in our case.

5.1.8 Heterogeneous Graphs

The basic structure of the KD-graph used in our model and other existing suggestion model is one type of heterogeneous graph that consists of multiple types of nodes and edges. There exist some research focus on the similarity search in heterogeneous graphs. PathRank [48] extends the Personalized PageRank algorithm on heterogeneous graphs by discriminating different paths during the random walk process guided by predefined meta-paths. Sharing the similar idea, the construction of the KD-graph widely used in the query suggestion is actually guided by so-call meta-path KEYWORD-DOCUMENT-KEYWORD. However, the PathRank does not provide a specific edge-weighing method suitable for our problem. PathSim [49] and PCRW [50] measure the proximity between nodes in heterogeneous graphs by only exploiting the given meta-path without considering the overall structure of given heterogeneous graph, which is inferior to PathRank. In addition, this line of research pay more attention to the graph models and apply existing algorithms for computing the most similar nodes, while algorithm PA proposed in this paper is also applicable to clusterable heterogeneous graphs for the purpose of computation acceleration.

5.2 Random Walk Computation

Random walk with restart, also known as Personalized PageRank, has been widely used for node similarity measures in graph data, especially since its successful application by the Google search engine [51].

5.2.1 Pre-Computation Based Approaches

Matrix-based methods [20], [21] solve PPR by pre-computing the inversion matrix. Tong et al. [20] propose a matrix-based approach B_LIN that reduces the pre-computation cost of the full matrix inversion by partitioning the graph. Fujiwara et al. [21] propose a K-dash method that finds the top- k nodes with the highest PPR scores, based on a LU decomposition of the transition matrix. Alternative to matrix-based approaches, Monte Carlo (MC) methods [22], [23], [24] can simulate the RWR process. Fogaras et al. [22] propose to approximate PPR by pre-computing and approximating for each node u a set of ending vertices for random walks starting from u . If u later becomes a query node, its PPR is approximated according to the distribution of these vertices. Similarly, Bahmani et al. [23] approximate PPR by counting the number of times a node is visited by pre-computed random walks. All above methods require the apriori knowledge of the complete graph; however, in our problem, the edge weights are dynamically adjusted according to the user location, thus these approaches are inapplicable.

5.2.2 Online Approaches

MC can also be applied online, without relying on pre-computations; a number of random walks are tried from the query node and the PPR score of other nodes are estimated from these samples [24]. However, as shown later in [28], a large number of (expensive) random walks are required in order to achieve acceptable precision. Fujiwara et al. [28] propose a method for efficient ad-hoc top- k PPR search with exact node ranking. They compute the random walk (without restart) probabilities of the nodes, and employ the probabilities to estimate upper/lower bounds of the candidate nodes. This approach is applicable when the complete transition matrix is available beforehand, however, obtaining the complete transition matrix in our problem involves the multiplication of two matrices and it is very expensive. Berkhin [25] proposes the Bookmark Coloring Algorithm to derive an approximation of the PPR vector. Gupta et al. [27] extend BCA and employ early termination heuristics for the top- k PPR calculation. We extend BCA [25] as our baseline algorithm for location-aware suggestions.

6 CONCLUSION

In this paper, we proposed an LKS framework providing keyword suggestions that are relevant to the user information needs and at the same time can retrieve relevant documents near the user location. A baseline algorithm extended from algorithm BCA [25] is introduced to solve the problem. Then, we proposed a partition-based algorithm which computes the scores of the candidate keyword queries at the partition level and utilizes a lazy mechanism to greatly reduce the computational cost. Empirical studies are conducted to study the effectiveness of our LKS framework

and the performance of the proposed algorithms. The result shows that the framework can offer useful suggestions and that PA outperforms the baseline algorithm significantly. In the future, we plan to further study the effectiveness of the LKS framework by collecting more data and designing a benchmark. In addition, subject to the availability of data, we will adapt and test LKS for the case where the locations of the query issuers are available in the query log. Finally, we believe that PA can also be applied to accelerate RWR on general graphs with dynamic edge weights; we will investigate this potential in the future.

ACKNOWLEDGMENTS

This work was funded by EC grant 657347/H2020-MSCA-IF-2014 and by GRF grant 17205015 from Hong Kong RGC. Dingming Wu is the corresponding author.

REFERENCES

- [1] R. Baeza-Yates, C. Hurtado, and M. Mendoza, "Query recommendation using query logs in search engines," in *Proc. Int. Conf. Current Trends Database Technol.*, 2004, pp. 588–596.
- [2] D. Beeferman and A. Berger, "Agglomerative clustering of a search engine query log," in *Proc. 6th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2000, pp. 407–416.
- [3] H. Cao, D. Jiang, J. Pei, Q. He, Z. Liao, E. Chen, and H. Li, "Context-aware query suggestion by mining click-through and session data," in *Proc. 14th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2008, pp. 875–883.
- [4] N. Craswell and M. Szummer, "Random walks on the click graph," in *Proc. 30th Annu. Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2007, pp. 239–246.
- [5] Q. Mei, D. Zhou, and K. Church, "Query suggestion using hitting time," in *Proc. 17th ACM Conf. Inf. Knowl. Manage.*, 2008, pp. 469–478.
- [6] Y. Song and L.-W. He, "Optimal rare query suggestion with implicit user feedback," in *Proc. 19th Int. Conf. World Wide Web*, 2010, pp. 901–910.
- [7] T. Miyanishi and T. Sakai, "Time-aware structured query suggestion," in *Proc. 36th Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2013, pp. 809–812.
- [8] A. Anagnostopoulos, L. Becchetti, C. Castillo, and A. Gionis, "An optimization framework for query recommendation," in *Proc. ACM Int. Conf. Web Search Data Mining*, 2010, pp. 161–170.
- [9] P. Boldi, F. Bonchi, C. Castillo, D. Donato, A. Gionis, and S. Vigna, "The query-flow graph: Model and applications," in *Proc. 17th ACM Conf. Inf. Knowl. Manage.*, 2008, pp. 609–618.
- [10] Y. Song, D. Zhou, and L.-w. He, "Query suggestion by constructing term-transition graphs," in *Proc. 5th ACM Int. Conf. Web Search Data Mining*, 2012, pp. 353–362.
- [11] L. Li, G. Xu, Z. Yang, P. Dolog, Y. Zhang, and M. Kitsuregawa, "An efficient approach to suggesting topically related web queries using hidden topic model," *World Wide Web*, vol. 16, pp. 273–297, 2013.
- [12] D. Wu, M. L. Yiu, and C. S. Jensen, "Moving spatial keyword queries: Formulation, methods, and analysis," *ACM Trans. Database Syst.*, vol. 38, no. 1, pp. 7:1–7:47, 2013.
- [13] D. Wu, G. Cong, and C. S. Jensen, "A framework for efficient spatial web object retrieval," *VLDB J.*, vol. 21, no. 6, pp. 797–822, 2012.
- [14] J. Fan, G. Li, L. Zhou, S. Chen, and J. Hu, "SEAL: Spatio-textual similarity search," *Proc. VLDB Endowment*, vol. 5, no. 9, pp. 824–835, 2012.
- [15] P. Bouros, S. Ge, and N. Mamoulis, "Spatio-textual similarity joins," *Proc. VLDB Endowment*, vol. 6, no. 1, pp. 1–12, 2012.
- [16] Y. Lu, J. Lu, G. Cong, W. Wu, and C. Shahabi, "Efficient algorithms and cost models for reverse spatial-keyword k -nearest neighbor search," *ACM Trans. Database Syst.*, vol. 39, no. 2, pp. 13:1–13:46, 2014.
- [17] S. Basu Roy and K. Chakrabarti, "Location-aware type ahead search on spatial databases: Semantics and efficiency," in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2011, pp. 361–372.
- [18] R. Zhong, J. Fan, G. Li, K.-L. Tan, and L. Zhou, "Location-aware instant search," in *Proc. 21st ACM Conf. Inf. Knowl. Manage.*, 2012, pp. 385–394.
- [19] I. Miliou and A. Vlachou, "Location-aware tag recommendations for flickr," in *Proc. 25th Int. Conf. Database Expert Syst. Appl.*, 2014, pp. 97–104.
- [20] H. Tong, C. Faloutsos, and J.-Y. Pan, "Fast random walk with restart and its applications," in *Proc. 6th Int. Conf. Data Mining*, 2006, pp. 613–622.
- [21] Y. Fujiwara, M. Nakatsuji, M. Onizuka, and M. Kitsuregawa, "Fast and exact top- k search for random walk with restart," *Proc. VLDB Endowment*, vol. 5, no. 5, pp. 442–453, Jan. 2012.
- [22] D. Fogaras, B. Rácz, K. Csalogány, and T. Sarlós, "Towards scaling fully personalized PageRank: Algorithms, lower bounds, and experiments," *Internet Math.*, vol. 2, no. 3, pp. 333–358, 2005.
- [23] B. Bahmani, A. Chowdhury, and A. Goel, "Fast incremental and personalized PageRank," *Proc. VLDB Endowment*, vol. 4, no. 3, pp. 173–184, Dec. 2010.
- [24] K. Avrachenkov, N. Litvak, D. Nemirovsky, E. Smirnova, and M. Sokol, "Quick detection of top- k personalized PageRank lists," in *Proc. 8th Int. Workshop Algorithms Models Web Graph*, 2011, vol. 6732, pp. 50–61.
- [25] P. Berkhin, "Bookmark-coloring algorithm for personalized pagerank computing," *Internet Math.*, vol. 3, pp. 41–62, 2006.
- [26] G. Jeh and J. Widom, "Scaling personalized web search," in *Proc. 12th Int. Conf. World Wide Web*, 2003, pp. 271–279.
- [27] M. Gupta, A. Pathak, and S. Chakrabarti, "Fast algorithms for top- k personalized pagerank queries," in *Proc. 17th Int. Conf. World Wide Web*, 2008, pp. 1225–1226.
- [28] Y. Fujiwara, M. Nakatsuji, H. Shiokawa, T. Mishima, and M. Onizuka, "Efficient ad-hoc search for personalized PageRank," in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2013, pp. 445–456.
- [29] I. S. Dhillon, "Co-clustering documents and words using bipartite spectral graph partitioning," in *Proc. ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2001, pp. 269–274.
- [30] G. Pass, A. Chowdhury, and C. Torgeson, "A picture of search," in *Proc. 1st Int. Conf. Scalable Inf. Syst.*, 2006.
- [31] S. Bhatia, D. Majumdar, and P. Mitra, "Query suggestions in the absence of query logs," in *Proc. Annu. Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2011, pp. 795–804.
- [32] L. Chen, G. Cong, C. S. Jensen, and D. Wu, "Spatial keyword query processing: an experimental evaluation," *Proc. VLDB Endowment*, vol. 6, pp. 217–228, 2013.
- [33] R. Baeza-Yates and A. Tiberi, "Extracting semantic relations from query logs," in *Proc. 13th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, 2007, pp. 76–85.
- [34] H. Ma, H. Yang, I. King, and M. R. Lyu, "Learning latent semantic relations from clickthrough data for query suggestion," in *Proc. 17th ACM Conf. Inf. Knowl. Manage.*, 2008, pp. 709–718.
- [35] H. Deng, I. King, and M. R. Lyu, "Entropy-biased models for query representation on the click graph," in *Proc. 32nd Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2009, pp. 339–346.
- [36] K. Järvelin and J. Kekäläinen, "Cumulated gain-based evaluation of ir techniques," *ACM Trans. Inf. Syst.*, vol. 20, no. 4, pp. 422–446, Oct. 2002.
- [37] B. Sarwar, G. Karypis, J. Konstan, and J. Riedl, "Item-based collaborative filtering recommendation algorithms," in *Proc. 10th Int. Conf. World Wide Web*, 2001, pp. 285–295.
- [38] Y. Song, D. Zhou, and L.-w. He, "Post-ranking query suggestion by diversifying search results," in *Proc. 34th Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2011, pp. 815–824.
- [39] X. Zhu, J. Guo, X. Cheng, P. Du, and H.-W. Shen, "A unified framework for recommending diverse and relevant queries," in *Proc. 20th Int. Conf. World Wide Web*, 2011, pp. 37–46.
- [40] U. Ozertem, O. Chapelle, P. Donmez, and E. Velpasaoglu, "Learning to suggest: A machine learning framework for ranking query suggestions," in *Proc. 35th Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2012, pp. 25–34.
- [41] Y. Liu, R. Song, Y. Chen, J.-Y. Nie, and J.-R. Wen, "Adaptive query suggestion for difficult queries," in *Proc. 35th Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2012, pp. 15–24.
- [42] J.-R. Wen, J.-Y. Nie, and H.-J. Zhang, "Clustering user queries of a search engine," in *Proc. 10th Int. Conf. World Wide Web*, 2001, pp. 162–168.
- [43] R. Li, B. Kao, B. Bi, R. Cheng, and E. Lo, "DQR: A probabilistic approach to diversified query recommendation," in *Proc. 21st ACM Conf. Inf. Knowl. Manage.*, 2012, pp. 16–25.

- [44] Z. Zhang and O. Nasraoui, "Mining search engine query logs for query recommendation," in *Proc. 15th Int. Conf. World Wide Web*, 2006, pp. 1039–1040.
- [45] S. Cucerzan and R. W. White, "Query suggestion based on user landing pages," in *Proc. 30th Annu. Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, 2007, pp. 875–876.
- [46] J. Myllymaki, D. Singleton, A. Cutter, M. Lewis, and S. Eblen, "Location based query suggestion," U.S. Patent 8 301 639, Oct. 30, 2012.
- [47] T. Gaasterland, "Cooperative answering through controlled query relaxation," *IEEE Expert*, vol. 12, no. 5, pp. 48–59, Sep. 1997.
- [48] S. Lee, S. Park, M. Kahng, and S.-G. Lee, "PathRank: A novel node ranking measure on a heterogeneous graph for recommender systems," in *Proc. 21st ACM Int. Conf. Inf. Knowl. Manage.*, 2012, pp. 1637–1641.
- [49] Y. Sun, J. Han, X. Yan, P. S. Yu, and T. Wu, "Pathsim: Meta path-based top-k similarity search in heterogeneous information networks," *Proc. VLDB Endowment*, vol. 4, pp. 992–1003, 2011.
- [50] N. Lao and W. W. Cohen, "Relational retrieval using a combination of path-constrained random walks," *Mach. Learn.*, vol. 81, no. 1, pp. 53–67, 2010.
- [51] S. Brin and L. Page, "The anatomy of a large-scale hypertextual web search engine," *Comput. Netw. ISDN Syst.*, vol. 30, no. 1–7, pp. 107–117, Apr. 1998.



Shuyao Qi received the bachelor's degree of engineering from the Department of Computer Science and Technology, Zhejiang University, in 2012. He is working toward the PhD degree at the Department of Computer Science, University of Hong Kong. His research focuses on query processing on spatial and textual data.



Dingming Wu received the bachelor's and the master's degrees in computer science from the Huazhong University of Science and Technology and Peking University in 2005 and 2008, respectively. She received the PhD degree in computer science from Aalborg University in 2011. She is a postdoc fellow at the University of Hong Kong. She was a postdoc teaching fellow at Hong Kong Baptist University from 2011 to 2013. Her research concerns spatial keyword search, geo-social networks, and recommendation systems.



Nikos Mamoulis received the diploma in computer engineering and informatics from the University of Patras, Greece, and the PhD degree in computer science from the Hong Kong University of Science and Technology. Since 2001, he has been a faculty member in the Department of Computer Science, University of Hong Kong. Since 2014, he moved to the Department of Computer Science and Engineering, University of Ioannina. His research focuses on managing complex data types, privacy and security in databases, and uncertain data management.

▷ For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.